

# Literature Review: Deep Learning for Seizure Detection and Prediction from Non-EEG Wearable Signals (2013–2026)

Paulin Saher  
Universität zu Köln

January 23, 2026

## Contents

|          |                                                                            |           |
|----------|----------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                                        | <b>3</b>  |
| 1.1      | Problem Statement and Motivation . . . . .                                 | 3         |
| 1.2      | Key Research Question . . . . .                                            | 3         |
| <b>2</b> | <b>Theoretical Foundations for Ambulatory Seizure Monitoring</b>           | <b>3</b>  |
| 2.1      | The Ambulatory Monitoring Paradigm: Requirements and Challenges . . .      | 3         |
| 2.2      | Physiological Biosignals for Wearable Sensing . . . . .                    | 4         |
| 2.3      | Seizure Manifestations and Corresponding Biosignals . . . . .              | 4         |
| 2.4      | Evaluation Frameworks: From Retrospective Analysis to Prospective Validity | 4         |
| <b>3</b> | <b>Methodological Approach and Structure</b>                               | <b>5</b>  |
| <b>4</b> | <b>Search Strategy</b>                                                     | <b>7</b>  |
| 4.1      | PICO-Based Search Methodology . . . . .                                    | 7         |
| 4.2      | Exploration . . . . .                                                      | 8         |
| 4.3      | Practical Filters and Reproducibility . . . . .                            | 8         |
| <b>5</b> | <b>Results</b>                                                             | <b>14</b> |
| 5.1      | Study Characteristics . . . . .                                            | 14        |
| 5.2      | Modality Performance . . . . .                                             | 15        |
| 5.3      | Architecture and Algorithm Patterns . . . . .                              | 16        |
| 5.4      | Modalities, Architectures, and Personalization . . . . .                   | 18        |
| 5.5      | Personalization Strategies . . . . .                                       | 19        |
| 5.6      | Performance Metrics . . . . .                                              | 21        |
| 5.7      | Detection Versus Forecasting . . . . .                                     | 22        |
| 5.8      | Clinical Readiness and Deployment . . . . .                                | 24        |
| 5.9      | Detection, Forecasting, and Clinical Readiness . . . . .                   | 26        |

|          |                                             |           |
|----------|---------------------------------------------|-----------|
| <b>6</b> | <b>Discussion</b>                           | <b>28</b> |
| 6.1      | Clinical Implications . . . . .             | 28        |
| 6.2      | Technical Insights and Trade-offs . . . . . | 29        |
| 6.3      | Limitations of the Evidence Base . . . . .  | 31        |
| 6.4      | Future Directions . . . . .                 | 32        |
| <b>7</b> | <b>Conclusion</b>                           | <b>33</b> |

# Abstract

## Background

Wearable devices using non-EEG biosignals offer potential for continuous ambulatory seizure monitoring without the burden of scalp electrodes. This systematic review examines the current state of deep learning approaches for seizure detection and forecasting using accelerometry, cardiac signals, and other wearable modalities.

## Methods

Following PRISMA guidelines, 13 primary studies (9 detection, 4 forecasting) published between 2013 and 2026 were reviewed. Studies were evaluated following the framework proposed by Webster and Watson, 2002 based on the following metrics: sensitivity, false alarm rate (FAR), validation methodology, and ILAE phase of development.

## Results

Detection of convulsive seizures achieves high sensitivity (71.6-100%) using primarily accelerometry-based approaches. Only two systems approach the ILAE Phase 3 benchmark for validated devices: Fine et al. (tonic seizures, 100% sensitivity, 0.023/h FAR) and Dong et al. (nocturnal severe seizures, 71.6% sensitivity, 0.165/h FAR). ECG-based detection consistently fails clinical utility due to excessive FAR (1.91-39.75/h). Nighttime monitoring substantially outperforms daytime (1/61 nights vs. 1/9 days).

Seizure forecasting remains immature. When evaluated with rigorous leave-one-subject-out cross-validation, only 43.5% of patients achieve better-than-chance performance. Patient-specific tuning inflates performance by 30-40 percentage points, creating unrealistic expectations for real-world deployment. Circadian patterns provide the most reliable forecasting signals, but wearable-only features fail at daily resolution.

Validation rigor is insufficient: only 3 of 13 studies employ prospective designs, and leave-one-subject-out cross-validation remains rare. Eleven of 13 studies validate in controlled epilepsy monitoring unit settings, likely overestimating real-world performance.

## Conclusions

Wearable seizure detection for convulsive seizures has achieved clinical viability for specific use cases, particularly nocturnal monitoring. However, seizure forecasting faces a fundamental responder problem that prevents clinical translation. False alarm rates remain the primary barrier for daytime detection, and real-world prospective validation is urgently

needed. Future research must prioritize standardized metrics, edge deployment optimization, and identification of baseline biomarkers to predict which patients will benefit from wearable monitoring systems.

# 1 Introduction

## 1.1 Problem Statement and Motivation

Epilepsy, a chronic neurological disorder affecting 60 million people worldwide, with approximately one-third of patients remaining drug-resistant (Hussein et al., 2018). Sveinsson et al., 2020 found that 69% of Deaths due to a Generalized tonic-clonic seizure could have been prevented if not left unattended, highlighting the demand for reliable and timely prediction.

While reliable prediction is already possible through electroencephalography (EEG), its operational complexity render it unsuitable for ambulatory use. To address the need for ambulatory detection many are experimenting with utilizing autonomic and motion signals to detect and predict seizures.

Translating and analysing these often noisy and complex time-series physiological signals has pushed traditional machine learning (ML) approaches to their limits. These approaches require manual engineering such as cleaning and extracting of the data, often failing to capture the intricate and non-linear complex temporal dynamic of these signals. It is this limitation that deep learning (DL) models are designed to overcome, with their ability of hierarchical and learned feature extraction uniquely addressing these limitations. This review therefore examines and assesses recent advances in time-series deep learning - particularly convolutional neural networks (CNNs) and long short term memory (LSTMs) - applied to seizure detection and prediction using non-EEG wearable biosignals. Central to this analysis is the question of how architectural choices, training paradigms (e.g., transfer learning, personalization), and thorough prospective evaluation can converge to meet the stringent clinical requirements of high sensitivity and minimal false alarm rates (FAR) in real-world deployment.

## 1.2 Key Research Question

How do different deep learning architectures and biosignal modalities excluding EEG compare in their ability to achieve a optimal trade-off between sensitivity and false alarm rate for ambulatory seizure monitoring?

# 2 Theoretical Foundations for Ambulatory Seizure Monitoring

## 2.1 The Ambulatory Monitoring Paradigm: Requirements and Challenges

Beniczky et al., 2021, Chen et al., 2022 found that translation of seizure detection from the controlled clinical environment to an ambulatory setting redefines approaches and performance requirements. In clinical environments maximum sensitivity is prioritized over FAR, when engineering for ambulatory use a balance of both metrics is needed.

Failing to detect a seizure (low sensitivity) poses a safety risk, While excessive false alarms (high FAR) leads to desensitizing and "alarm fatigue" and the device being ignored or not used at all. Users reported that devices should achieve high sensitivity ( $\geq 90\%$ ) with low false alarm rates — acceptable FARs were 1-2/week for patients with  $\geq 1$  seizure/week and 1-2/month for patients with  $< 1$  seizure/week (Sivathamboo et al., 2022). While more complex models can find more intricate patterns, they also introduce the "black-box" problem wherein the model lacks transparency AbuAlrob et al., 2025; Ghaderi, 2025. The design should incorporate explanations for decisions to ensure user trust and regulatory compliance AbuAlrob et al., 2025; Miron et al., 2025.

## 2.2 Physiological Biosignals for Wearable Sensing

Wearable devices circumvent the impracticality of EEG by capturing peripheral biosignals that serve as proxies for central nervous system activity during seizures. These proxies will inherently relay incomplete and noisy signals so a combination of multiple sensors is required. The acceleration/attitude sensor excels at detecting muscle convulsions but is noisy in daily use, as it captures all movement. The electrodermal activity (EDA) sensor detects autonomic arousal by measuring the skin conductance associated with sweat gland activity.

## 2.3 Seizure Manifestations and Corresponding Biosignals

Wearable biosignals act as peripheral proxies for cerebral activity during seizures. Different seizure types produce distinct physiological signatures across these signals, informing detection approaches. This relationship centers on two primary manifestations: motor activity and autonomic activation.

Convulsive seizures, such as generalized tonic-clonic events, produce clear motor signatures captured by accelerometry (ACC), which measures movement and posture changes characteristic of ictal events (Wang et al., 2025; Wu et al., 2024).

At the same time, seizures drive autonomic responses recorded through complementary biosignals. Electrodermal activity (EDA) measures sympathetic nervous system arousal via skin conductance changes (Jain et al., 2017; Miron et al., 2025). Cardiac activity provides particularly strong autonomic proxies: electrocardiography (ECG) captures heart rate and rhythm alterations (Ghaderi, 2025; Leal et al., 2017), while heart rate variability (HRV) quantifies autonomic balance dynamics (Mason et al., 2024; Pavei et al., 2017).

In wearable applications, photoplethysmography (PPG) serves as a practical alternative for cardiac monitoring, deriving pulse rate and variability though with greater motion sensitivity than ECG (Chen et al., 2022; Seth et al., 2023).

As a result, detection strategies combine these biosignals according to target seizure types, with multimodal fusion providing the most reliable approach for ambulatory monitoring.

## 2.4 Evaluation Frameworks: From Retrospective Analysis to Prospective Validity

The clinical validity of a seizure detection algorithm is determined by the rigor of its evaluation, formalized in the ILAE’s phased framework (Phases 0-4) (Beniczky & Ryvlin, 2018; Beniczky et al., 2021). This framework classifies studies into five phases: Phase 0 (proof of concept), Phase 1 (retrospective, small sample), Phase 2 (minimum 10 patients

with seizures, 15 recorded seizures), Phase 3 (prospective, multicenter,  $\geq 30$  seizures from  $\geq 20$  patients, real-time detection), and Phase 4 (real-world home validation).

Retrospective studies (Phases 0-2), while foundational, risk substantial data leakage and optimistic bias if they use non-chronological data splits, allowing future information to inform training (Kalousios et al., 2024; Wong et al., 2023). Pseudo-prospective validation addresses this by enforcing chronological order, training only on past seizures to test on subsequent ones, providing a more realistic performance estimate (Kalousios et al., 2024; Lu et al., 2025).

The clinical gold standard is a prospective trial (Phase 3), which requires a locked algorithm, a video-EEG reference standard, and testing on new, unseen data from multiple centers. The ILAE systematic review identified only three Phase 3 studies meeting these strict validation standards (Beniczky et al., 2021). These demonstrated 90-96% sensitivity for generalized tonic-clonic seizures (GTCS) and focal-to-bilateral tonic-clonic seizures (FBTCS) with false alarm rates of 0.2-0.67 per 24 hours (approximately 0.008-0.028/h). This represents the clinically validated performance benchmark for convulsive seizure detection.

What constitutes an acceptable false alarm rate depends on clinical context and user perspective. Systematic surveys reveal important differences in tolerance: patients and caregivers typically require 100% sensitivity and accept approximately one false alarm per seizure (or one per week for seizure-free patients). In contrast, clinicians consider 90% sensitivity adequate and tolerate between two false alarms per week and one per month, depending on seizure frequency. This gap between user and clinician expectations creates a fundamental tension for device development.

The final step is in-field evaluation (Phase 4) in patients' homes, where metrics like the false alarm rate per hour (FAR/h) become decisive for practical usability, as demonstrated by long-term studies like Nightwatch (Miron et al., 2025; Shum & Friedman, 2021).

Currently, the ILAE recommends wearable devices only for GTCS/FBTCS detection in unsupervised patients where alarms can enable rapid intervention - a weak/conditional recommendation based on moderate evidence. For all other seizure types, the ILAE does not recommend clinical use of currently available wearable devices (Beniczky et al., 2021).

### 3 Methodological Approach and Structure

The review follows structured IS/medical literature guidance (concept matrix inspired by (Webster & Watson, 2002) and PRISMA principles (Tugwell & Tovey, 2021)). The methodological approach uses a multi-stage search and transparent extraction and synthesis procedures as outlined below.

#### 1. Search strategy (Multi-stage, 2015-2025):

- **Scoping:** Google Scholar for broad coverage and identifying candidate papers.
- **Formal queries:** IEEE Xplore, PubMed, and Scopus using controlled keyword strings and deduplication.
- **Citation chasing:** Backward and forward citation tracking on key papers to ensure completeness.

#### 2. Inclusion / Exclusion criteria:

- **Inclusion:** Studies using wearable/non-EEG autonomic signals for seizure detection or preictal prediction.
- **Exclusion:** EEG-only studies, invasive recordings, or papers limited to purely algorithmic simulations without human data.

Extracted dimensions populate two concept matrix tables with the following fields:

**1. Study Matrix (Table 1 for detection, Table 2 for forecasting):**

- Study identification (author, year)
- Objective (detection vs forecasting, seizure type)
- Design (prospective, retrospective, pseudo-prospective, LOSO CV)
- Sample (participant count, seizure count)
- Device and sensor location
- Algorithm architecture
- Performance metrics (sensitivity, specificity, FPR)
- Key limitations

**2. Detailed Metrics Matrix (Table 3 for detection, Table 4 for forecasting):**

- Validation approach (independent test, cross-validation, LOSO)
- Detection latency
- Patient-level success rate
- Precision/PPV
- Additional clinical metrics (HMS, IoC, AUC, F1)
- Clinical notes

**Additional Extracted Dimensions:** The following dimensions were extracted and analyzed to inform the narrative synthesis and dimensional analysis presented in Section 5, complementing the tabular comparison:

- Setting (clinical vs ambulatory)
- Modality details (ECG/HRV, PPG, EDA, ACC)
- Evaluation granularity (sample- vs alarm-based)
- Windowing specifics (window length, stride, overlap)
- Personalization strategies
- Interpretability and safety considerations
- Data characteristics and dataset identity
- Missing data handling (imputation, exclusion) and synthetic data use
- Inference requirements (latency, computational load, hardware)

This two-table approach follows Webster & Watson’s concept matrix methodology (Webster & Watson, 2002), balancing comprehensive data extraction with readable summary tables while enabling detailed dimensional analysis in the narrative synthesis.

**Prioritization Criteria:** Studies including splits preserving patient-specific data, FAR/h reporting and external or cross validation will be prioritized in the synthesis. A PRISMA-style flow diagram will summarize and visualize screening and inclusion.

## 4 Search Strategy

### 4.1 PICO-Based Search Methodology

The literature search was structured using the PICO framework (Patient/Problem, Intervention, Comparison, Outcome) to identify relevant studies on deep learning for non-EEG epilepsy detection. The following key terms were used in combination with Boolean operators (AND, OR, NOT):

1. **P (Patient/Problem):** Epileptic seizures in need of detection or prediction.
  - **Keywords:** ‘epilepsy’, ‘seizure’, ‘ictal’, ‘preictal’
  - **MeSH Terms (PubMed):** Epilepsy, Seizures
2. **I (Intervention):** The diagnostic method based on wearable biosignals and deep learning.
  - **Biosignal Keywords:** ‘electrocardiogram’, ‘ECG’, ‘heart rate variability’, ‘HRV’, ‘photoplethysmography’, ‘PPG’, ‘electrodermal activity’, ‘EDA’, ‘accelerometer’, ‘ACC’
  - **Technology Keywords:** ‘wearable’, ‘wearable device’
  - **Method Keywords:** ‘deep learning’, ‘machine learning’, ‘convolutional neural network’, ‘CNN’, ‘long short-term memory’, ‘LSTM’, ‘neural network’
  - **MeSH Terms (PubMed):** Wearable Electronic Devices, Electrocardiography, Deep Learning
3. **C (Comparison):** The explicit exclusion of studies focused solely on electroencephalography (EEG) as the primary modality.
  - **Exclusion Keywords:** ‘EEG’, ‘electroencephalography’, ‘iEEG’, ‘intracranial’
  - **Search Logic:** The NOT operator was applied to this group in all database searches.
4. **O (Outcome):** Metrics quantifying the performance of a detection or prediction system.
  - **Keywords:** ‘detection’, ‘prediction’, ‘sensitivity’, ‘specificity’, ‘false alarm rate’, ‘FAR’, ‘area under the curve’, ‘AUC’

## 4.2 Exploration

Due to the insufficient amount of papers gathered from systematic review an exploration was done on the already existing papers. The exploration consisted of graph tools to visualize backlinks and gemini deep research feature to find relevant papers. Through this method an additional 4 papers were acquired bringing the total number of papers up to 13 and expanding the time frame to 2013-2026.

## 4.3 Practical Filters and Reproducibility

Consistent with the overall review scope, searches were restricted to **2015-2025** and **journal articles** with few exceptions made regarding **conference proceedings**. All database searches were documented (query string, platform, and execution date). Results were deduplicated prior to screening.

**Note on exclusions:** Because some non-EEG wearable studies rely on video-EEG for labeling, EEG-related terms were used to filter **EEG-only modalities**, not to exclude studies that mention EEG solely as a reference standard.




Figure 1: PRISMA flow diagram of the study selection process.

Table 1: Webster &amp; Watson Concept Matrix: Detection Studies

| Study                       | Design      | Sample         | Device                 | Loc.      | Para. | Mod. | Pers. | Algorithm                                  | Sens                   | Spec                  | FPR          |
|-----------------------------|-------------|----------------|------------------------|-----------|-------|------|-------|--------------------------------------------|------------------------|-----------------------|--------------|
| <b>Spahr et al. 2025</b>    | Prospective | n=384, 49 CSs  | Empatica E4 (ACC)      | Wrist     | Ens   | 1    | Glb   | Ensemble 1D CNN (30 models)                | 96%                    | –                     | 0.0054/h     |
| <b>Reintjes et al. 2025</b> | Retro       | n=120, 856 szs | Single-lead ECG        | Chest     | Anom  | 1    | Subj  | Anomaly: Matrix Profile, MADRID, TimeVQVAE | 2.44–82.79%            | –                     | 0.11–65.62/h |
| <b>Fine 2025</b>            | Phase 1     | n=15/3         | 6-axis band (ACC+Gyro) | Wrist     | Feat  | M    | Glb   | ANN (594 hand-crafted feat)                | 100%Test, 96% complete | –                     | 0.023/h      |
| <b>Dong et al. 2026</b>     | Prospective | n=68, 1846 szs | NightWatch             | Arm       | 2-stg | M    | Glb   | CNN-LSTM + Attention                       | 71.6%                  | –                     | 0.165/h      |
| <b>Wang et al. 2025</b>     | –           | n=28, 62 szs   | Biovital-P1            | Wrist     | Feat  | M    | Glb   | LSTM (40 hidden)                           | 56.40% to 95.3%        | –                     | 0.364/h      |
| <b>Singh Rathore 2024</b>   | Case study  | 1 pt, 6 hrs    | Wearable (EDA+ACC+HR)  | –         | Feat  | M    | CS    | MLP (multi-modal features)                 | 96.8%                  | 94.8% <sup>prec</sup> | –            |
| <b>Elemam et al. 2025</b>   | Cross-sect  | Q:198, HRV:30  | Camera PPG             | Thumbs    | Hyb   | M    | Glb   | CNN + Rule-based fusion                    | 95.1% Q, 93% HRV       | 97.06% Q              | –            |
| <b>Borujeny 2013</b>        | 3-pt        | 3 pts, 20 szs  | MICAz ACC              | Arm/thigh | Feat  | 1    | Glb   | ANN (time-domain feat)                     | 85%                    | –                     | 3 FA         |

**Note:** ACC = accelerometer, ANN = artificial neural network, CNN = convolutional neural network, CS = convulsive seizure / case study, ECG = electrocardiogram, EDA = electrodermal activity, HR = heart rate, HRV = HR variability, LSTM = long short-term memory, PPG = photoplethysmography, Sens = sensitivity, Spec = specificity, sz(s) = seizure(s), FPR = false positive rate, – = not reported. Para. = learning paradigm (Feat = feature-based, Ens = ensemble, Anom = anomaly detection, 2-stg = two-stage, Hyb = hybrid). Mod. = modality count (1 = single-modal, M = multi-modal). Pers. = personalization (Glb = global model, Subj = subject split, CS = case study). <sup>prec</sup> = precision (not specificity).

Table 2: Webster &amp; Watson Concept Matrix: Forecasting Studies

| Study              | Design        | Sample        | Device      | Loc.        | Para. | Mod. | Pers. | Algorithm                       | Sens               | Spec | FPR            |
|--------------------|---------------|---------------|-------------|-------------|-------|------|-------|---------------------------------|--------------------|------|----------------|
| Vieluf et al. 2025 | Retro         | n=70, 5437 d  | Embrace     | Wrist       | Hyb   | M    | Mix   | DNN + harmonic features + diary | 82%                | 67%  | –              |
| Meisel et al. 2020 | LOSO CV       | n=69, 452 szs | Empatica E4 | Wrist/Ankle | E2E   | M    | Glb   | LSTM (10 units)                 | 51.2%              | –    | TiW 43.7%      |
| Stirling 2021      | Retro+pseudon | n=11, 136 szs | Fitbit      | Wrist       | Feat  | M    | Pat   | LSTM+RF+LR ensemble (HRV feat)  | –                  | –    | AUC 0.74       |
| Nasseri 2021       | Retro         | n=6           | Empatica E4 | Wrist       | E2E   | M    | Tmp   | LSTM 4-layer (128 hidden)       | AUC 0.75 (SD 0.15) | –    | TiW 0.9–7.2h/d |
| Ode et al. 2023    | Retro         | n=66, 85 szs  | ECG         | –           | Anom  | 1    | Pat   | Self-Attentive AE               | 74%                | –    | 0.85/h         |

**Note:** ACC = accelerometer, AE = autoencoder, CNN = convolutional neural network, DNN = deep neural network, ECG = electrocardiogram, EDA = electrodermal activity, LSTM = long short-term memory, LOSO = leave-one-subject-out, PPG = photoplethysmography, Sens = sensitivity, Spec = specificity, sz(s) = seizure(s), FPR = false positive rate, TiW = time in warning, AUC = area under curve, – = not reported. Para. = learning paradigm (Feat = feature-based, E2E = end-to-end, Anom = anomaly detection, Hyb = hybrid). Mod. = modality count (1 = single-modal, M = multi-modal). Pers. = personalization (Glb = global model, Pat = patient-specific, Mix = mixed, Tmp = temporal split).

Table 3: Architecture and Modality Deep-Dive: Detection Studies

| Study                       | Paradigm      | Arch. Details                              | Modality tails                             | De- Fusion | Personalization    | Trade-off fig               | Con- Deployment                 |
|-----------------------------|---------------|--------------------------------------------|--------------------------------------------|------------|--------------------|-----------------------------|---------------------------------|
| <b>Spahr et al. 2025</b>    | Ensemble      | 30 1D CNN models, 14 conv layers, quantile | ACC 32 Hz, Euclidean norm, 30 s            | –          | Global, tunable    | 86–100%<br>0.0054/h         | @ Real-time (112 ms), on-device |
| <b>Reintjes et al. 2025</b> | Anomaly       | Matrix-Profile, MADRID, TimeVQVAE-AD       | ECG 256→8 Hz, bandpass                     | –          | Subject split      | 2.44–82.79%<br>0.11–65.62/h | @ Offline, hospital             |
| <b>Fine 2025</b>            | Feature-based | ANN, 594 hand-crafted features             | ACC+Gyro axis, 10 s intervals              | 6- Early   | Global, hold-out   | 100% @ 0.023/h              | Offline PC, EMU                 |
| <b>Dong et al. 2026</b>     | Two-stage     | Pre-screening + CNN-LSTM Attention         | + ACM 11-12→20 Hz, PPG 100→20 Hz, 5 min    | Early      | Global, 10-fold CV | 71.6% @ 0.165/h             | Real-time, Night-Watch, home    |
| <b>Wang et al. 2025</b>     | Feature-based | LSTM (40 hidden) + ReLU + FC               | ACC/GYRO 50 Hz, SEMG 200 Hz, EDA 4 Hz, 4 s | Early      | Global, hold-out   | 56.4–95.3%<br>0.364/h       | @ Real-time, hospital           |
| <b>Singh 2024</b>           | Feature-based | MLP, multi-modal features                  | EDA, HR/HRV, points                        | ACC, 25k   | Early              | Single pt (CS) 96.8% @ prec | 94.8% Real-time, cloud?         |
| <b>Elemam 2025</b>          | CNN-based     | CNN for HRV, PPG (separate)                | CNN for Audio 250 Hz, 10 s                 | No fusion  | Global, unclear    | 95.1% @ spec                | 97.06% Real-time, hospital      |
| <b>Borujeny 2013</b>        | Feature-based | ANN, time-domain features                  | ACC 2D 3 Hz, 9 s                           | –          | Global, 3 pts      | 85% @ 3 FA                  | Real-time, server (316 mW)      |

**Note:** Paradigm: Feature-based = handcrafted features, Ensemble = multiple models, Two-stage = sequential processing, Anomaly = unsupervised detection, CNN-based = convolutional approach. Arch. Details: architecture-specific components. Fusion: – = single-modal, Early = feature-level, No fusion = parallel separate models. Personalization: Global = population model, Subject split = train/test separated by subject. Trade-off Config: operating points with sensitivity @ FAR. Deployment: real-time capability, processing location, setting. ACC = accelerometer, ECG = electrocardiogram, PPG = photoplethysmography, EDA = electrodermal activity, HR = heart rate, HRV = HR variability, SEMG = surface electromyography, EMU = epilepsy monitoring unit.

Table 4: Architecture and Modality Deep-Dive: Forecasting Studies

| Study                     | Paradigm      | Arch. Details                           | Modality Details                        | Fusion        | Personalization           | Trade-off Config                    | Deployment                 |
|---------------------------|---------------|-----------------------------------------|-----------------------------------------|---------------|---------------------------|-------------------------------------|----------------------------|
| <b>Vieluf et al. 2025</b> | Hybrid        | DNN + harmonic modeling (24-h)          | EDA 4 Hz, ACC 32 Hz, Temp 1 Hz, 10 min  | Early         | Mixed, 5-fold + LOO       | 82% sens, 67% spec @ F1 0.81        | Offline MATLAB, home       |
| <b>Meisel et al. 2020</b> | End-to-end    | LSTM only (10 units)                    | EDA/ACC/BVP/Temp →4 Hz, 30 s            | Early         | Global, LOSO              | 51.2% @ IoC 14.1%                   | Offline, hospital          |
| <b>Stirling 2021</b>      | Feature-based | LSTM+RF+LR ensemble, LR combines        | HR 5 s, steps/sleep 1 min, hourly/daily | Decision      | Patient-specific, weekly  | AUC 0.74, 100% pts                  | Home (Fitbit), app         |
| <b>Nasseri 2021</b>       | End-to-end    | LSTM 4-layer, 128 hidden + FFT channels | ACC/BVP/EDA/Temp/HR →128 Hz, 60 s       | Early (17-ch) | Temporal split            | AUC 0.75 (SD 0.15) @ TiW 0.9–7.2h/d | Offline, cloud, home       |
| <b>Ode et al. 2023</b>    | Anomaly       | Self-Attentive AE (attention)           | ECG (RRI only), 45 s                    | –             | Patient-specific (99% CL) | 74% @ 0.85/h (99% CL)               | Real-time, cloud, hospital |

**Note:** Paradigm: Feature-based = handcrafted features, End-to-end = learns from raw data, Hybrid = combines both, Anomaly = unsupervised detection. Arch. Details: architecture-specific components. Fusion: – = single-modal, Early = feature-level, Decision = classifier output. Personalization: Global = population model, Patient-specific = individualized, Temporal = time-split within patient, LOSO = leave-one-subject-out, LOO = leave-one-out, Mixed = combination of approaches. Trade-off Config: operating points with sensitivity @ FAR. Deployment: real-time capability, processing location, setting. ACC = accelerometer, ECG = electrocardiogram, EDA = electrodermal activity, BVP = blood volume pulse, HR = heart rate, TEMP = temperature, FFT = fast Fourier transform, CL = confidence limit, AUC = area under curve, TiW = time in warning, IoC = Improvement over Chance.

Table 5: Performance Metrics Reporting Across Studies

| Metric                | Detection (n=9) | Forecasting (n=4) | Reported By                                                                             |
|-----------------------|-----------------|-------------------|-----------------------------------------------------------------------------------------|
| Sensitivity/Recall    | 8/8             | 1/4               | Spahr, Reintjes, Fine, Dong, Wang,<br>Singh Rathore, Elemam, Borujeny, Ode <sup>a</sup> |
| Specificity           | 3/9             | 0/4               | Elemam, Vieluf, Wang                                                                    |
| FPR/FA rate           | 8/9             | 0/4               | Spahr, Reintjes, Fine, Dong, Wang,<br>Singh Rathore, Elemam, Borujeny, Nasser, Ode      |
| AUC-ROC               | 3/9             | 3/4               | Dong, Reintjes, Ode, Nasser, Stirling                                                   |
| Detection Latency     | 4/9             | 1/4               | Spahr, Fine, Elemam, Borujeny, Nasser                                                   |
| Patient Success Rate  | 2/9             | 3/4               | Reintjes, Meisel, Stirling, Nasser, Ode                                                 |
| Precision/PPV         | 4/9             | 1/4               | Dong, Singh Rathore, Elemam, Ode                                                        |
| IoC / Time in Warning | 0/9             | 3/4               | Meisel, Stirling, Nasser                                                                |

## 5 Results

This systematic review analyzed 13 primary studies published between 2013 and 2026. The evidence base comprises 9 detection studies and 4 forecasting studies with a total of 912 participants. Table 1 and Table 2 present the study matrices for detection and forecasting approaches, respectively. Table 3 and Table 4 provide architectural details, and Table 5 summarizes the metrics reported across studies.

### 5.1 Study Characteristics

The 13 studies included in this review exhibit substantial variation in sample size and study design. Detection studies enrolled a median of 44 participants (range 1 to 384), while forecasting studies enrolled a median of 11 participants (range 6 to 70). The total sample comprises 687 participants in detection studies and 225 in forecasting studies.

Three studies employed prospective designs. Spahr et al. (2025) conducted a multi-center prospective study across eight epilepsy monitoring units with 384 patients. Dong et al. (2022) prospectively monitored 68 patients at home for up to three months. Fine (2025) prospectively recorded data from 15 patients in an epilepsy monitoring unit for tonic seizure detection. The remaining studies used retrospective analysis of existing datasets or small-scale feasibility designs.

Research activity has accelerated over time. Only one study was published between 2013 and 2015. Three forecasting studies were published between 2020 and 2022. Nine detection studies were published between 2023 and 2026. This pattern suggests increased research attention on detection approaches in recent years.

Eight studies were conducted in hospital or EMU settings. Six studies included home or ambulatory validation. Dong et al. (2022) achieved the longest home monitoring duration with 788 overnight recordings over three months. Stirling et al. (2021) achieved the longest overall monitoring period with a mean of 14.6 months per patient using consumer smartwatches.

**Key limitation.** The evidence base includes four studies with sample sizes below 20 participants. Singh Rathore et al. (2024) is a single-patient case study with limited generalizability. Borujeny et al. (2013) evaluated only 3 patients. Small sample sizes limit the reliability of performance estimates and generalizability to broader populations.

## 5.2 Modality Performance

Biosignal modality selection differs substantially between detection and forecasting approaches. Detection studies prioritize movement sensors, while forecasting studies emphasize autonomic nervous system indicators.

### 5.2.1 Modality Usage

Accelerometer data appears in nine of 13 studies (69%). Six detection studies and three forecasting studies use accelerometry. Electrodermal activity appears in five studies (38%). ECG or HRV data appears in five studies. PPG appears in four studies. Gyroscope data appears in three studies, all in detection applications.

Multi-modal approaches dominate the evidence base. Nine of 13 studies (69%) use multiple sensor types. Four studies use single-modality approaches: Spahr et al. (2025) (ACC only), Reintjes et al. (2025) (ECG only), Borujeny et al. (2013) (ACC only), and Ode et al. (2023) (ECG RRI only).

### 5.2.2 Sensor Location

Wrist-worn devices appear in eight studies (62%). Arm or armband placement appears in three studies. Chest placement appears in one study. Thigh placement appears in one study. Borujeny et al. (2013) used multiple sensor locations (right arm, left arm, left thigh).

The preference for wrist placement reflects patient acceptance and the availability of consumer wearable devices. The Empatica E4 wristband appears in three studies. The Embrace watch appears in one study. Fitbit smartwatches appear in one study.

### 5.2.3 Detection Performance by Modality

Single-modality accelerometer approaches achieve strong performance. Spahr et al. (2025) reported 96% sensitivity with 0.0054/h FPR using only ACC data from the Empatica E4. This represents the lowest FPR among all detection studies. Borujeny et al. (2013) reported 85% sensitivity using 2D ACC, though with only three patients.

Multi-modal movement detection results in excellent sensitivity in controlled settings. Fine (2025) used a 6-axis band combining ACC and gyroscope to achieve 100% sensitivity with 0.023/h FPR in an independent test set. However, the test set contained only 10 seizures from 3 patients.

ECG-based detection shows high FPR variability. Reintjes et al. (2025) compared three anomaly detection methods on single-lead ECG. Sensitivity ranged from 13.4% to 82.8%, while FPR ranged from 0.11/h to 65.62/h depending on method and configuration. The lowest FPR of 0.11/h detected only 2.44% of seizures. Ode et al. (2023) achieved 74% sensitivity with 0.85/h FPR using ECG RRI data.

Multi-modal approaches combining movement and autonomic signals show mixed results. Dong et al. (2022) combined ACC and PPG in a two-stage CNN-LSTM architecture, achieving 71.6% sensitivity with 0.165/h FPR. Wang et al. (2025) combined ACC, gyroscope, SEMG, and EDA to achieve 56.4% to 95.3% sensitivity depending on the feature set.

Wang et al. (2025) directly compared modality performance. Attitude angle features (PITCH and ROLL) outperformed raw accelerometer and gyroscope data for seizure

detection. This finding suggests that feature engineering can enhance the value of standard movement sensors.

#### 5.2.4 Forecasting Performance by Modality

Forecasting studies consistently use autonomic signals. Nasser et al. (2021) combined ACC, PPG, EDA, temperature, and HR to achieve AUC 0.75 (SD 0.15). Meisel et al. (2020) used EDA, PPG, temperature, and ACC from the Empatica E4 to achieve 51.2% sensitivity with IoC 14.1% in LOSO validation.

Vieluf et al. (2025) combined EDA, ACC, and temperature with diary data to achieve 82% sensitivity and 67% specificity. This study demonstrated that 24-hour harmonic patterns in EDA and accelerometry contain predictive information.

Stirling et al. (2021) focused on heart rate cycles combined with sleep and step data from Fitbit devices. The approach achieved AUC 0.74 with 100% patient success in hourly forecasting mode.

**Key finding.** Multi-modal approaches are common but not universally superior. Single-modality ACC (Spahr et al. (2025)) achieves the best FPR in the included detection literature. Forecasting consistently requires multi-modal autonomic data, but AUC plateaus around 0.75-0.8 regardless of modality combination.

### 5.3 Architecture and Algorithm Patterns

Deep learning architectures differ substantially between detection and forecasting applications. Detection studies employ diverse paradigms, while forecasting studies concentrate on LSTM-based approaches.

#### 5.3.1 Learning Paradigms

Feature-based approaches appear in four studies. These methods extract handcrafted features as input to classifiers. Fine (2025) extracted 594 features from 6-axis ACC and gyroscope data for an ANN classifier. Wang et al. (2025) engineered attitude angle features for an LSTM network. Singh Rathore et al. (2024) used multi-modal features for an MLP classifier.

End-to-end learning appears in two studies, both in forecasting. Meisel et al. (2020) used an LSTM with 10 units operating directly on raw multi-modal wrist data. Nasser et al. (2021) used a 4-layer LSTM with 128 hidden units receiving FFT channels of raw sensor data.

Hybrid approaches appear in two studies. Vieluf et al. (2025) combined a DNN with handcrafted harmonic features extracted from 24-hour patterns. Elemam et al. (2025) used separate CNNs for HRV and audio data without fusion between modalities.

Ensemble methods appear in detection applications. Spahr et al. (2025) trained 30 separate 1D CNN models with quantile aggregation. This ensemble approach achieved 96% sensitivity with 0.0054/h FPR. Dong et al. (2022) used a two-stage architecture with threshold-based pre-screening followed by a CNN-LSTM model.

Anomaly detection appears in two studies. Reintjes et al. (2025) compared three unsupervised anomaly detection methods on ECG data. Ode et al. (2023) used a self-attentive autoencoder for patient-specific anomaly detection on ECG RRI intervals.

### 5.3.2 Detection Architectures

CNN-based approaches dominate detection research. Spahr et al. (2025) implemented an ensemble of 30 CNN models with 14 convolutional layers each. The model operates on ACC data sampled at 32 Hz with 30-second windows. Elemam et al. (2025) used separate CNNs for HRV and audio processing.

LSTM appears in one detection study. Wang et al. (2025) used an LSTM with 40 hidden units processing attitude angle features. The architecture achieved 56.4% to 95.3% sensitivity depending on the feature configuration.

Traditional neural networks appear in earlier studies. Fine (2025) used an ANN with 594 handcrafted features. Borujeny et al. (2013) used an ANN with time-domain features from 2D ACC. Singh Rathore et al. (2024) used an MLP with multi-modal features.

### 5.3.3 Forecasting Architectures

LSTM architectures dominate forecasting research. Three of five forecasting studies use LSTM variants. Meisel et al. (2020) implemented a minimal LSTM with only 10 hidden units. Nasserri et al. (2021) used a deeper 4-layer LSTM with 128 hidden units per layer. Stirling et al. (2021) combined LSTM with random forest and logistic regression in an ensemble.

Non-LSTM deep learning appears in two studies. Vieluf et al. (2025) used a fully connected DNN with harmonic features. Ode et al. (2023) used a self-attentive autoencoder for anomaly detection.

### 5.3.4 Temporal Evolution

Architecture choices have evolved over time. The earliest study, Borujeny et al. (2013), used traditional ANN approaches. Studies published between 2020 and 2022 focused on LSTM-based forecasting architectures. Studies published between 2023 and 2026 employed more complex architectures including ensembles, hybrid methods, and attention mechanisms.

**Key finding.** Detection research employs architectural diversity including CNNs, LSTMs, ANNs, and anomaly detection. Forecasting research concentrates heavily on LSTM variants. No forecasting study uses CNN-based approaches, possibly because forecasting requires modeling temporal dependencies rather than spatial patterns.

### 5.3.5 Input Representations

Handcrafted features remain competitive. Three of eight detection studies use feature-based approaches with strong performance. Fine (2025) achieved 100% sensitivity using 594 handcrafted features. This suggests that domain knowledge remains valuable despite the end-to-end learning trend.

Window sizes vary substantially between applications. Detection studies use shorter windows ranging from 4 seconds to 5 minutes. Forecasting studies use longer windows ranging from 30 seconds to 60 minutes. Longer windows in forecasting may capture pre-ictal physiological changes.

**Limitation.** Architecture descriptions vary in detail across studies. Some studies report complete architectural details while others provide minimal information. This variability complicates direct comparison and replication.

## 5.4 Modalities, Architectures, and Personalization

### 5.4.1 Modality Selection and Sensor Placement

Biosignal modality selection differs between detection and forecasting approaches. Detection studies prioritize movement sensors, while forecasting studies emphasize autonomic nervous system indicators. Accelerometer data appears in nine of 13 studies (69%), with six detection studies and three forecasting studies using accelerometry. Electrodermal activity appears in five studies (38%). ECG or HRV data appears in five studies. PPG appears in four studies. Gyroscope data appears in three studies, all in detection applications.

Multi-modal approaches dominate the evidence base. Nine of 13 studies (69%) use multiple sensor types. Four studies use single-modality approaches: Spahr et al. (2025) (ACC only), Reintjes et al. (2025) (ECG only), Borujeny et al. (2013) (ACC only), and Ode et al. (2023) (ECG RRI only). Wrist-worn devices appear in eight studies (62%), reflecting patient acceptance and the availability of consumer wearable devices.

Single-modality accelerometer approaches achieve strong performance. Spahr et al. (2025) reported 96% sensitivity with 0.0054/h FPR using only ACC data, representing the lowest FPR among all detection studies. Multi-modal movement detection shows excellent sensitivity in controlled settings, with Fine (2025) achieving 100% sensitivity using a 6-axis band combining ACC and gyroscope. ECG-based detection shows high FPR variability, with Reintjes et al. (2025) reporting FPR ranging from 1.91/h to 65.62/h depending on the anomaly detection method.

Forecasting studies consistently use autonomic signals. Nasser et al. (2021) combined ACC, PPG, EDA, temperature, and HR to achieve AUC 0.75 (SD 0.15) across the full cohort, with significant predictors achieving mean AUC 0.80 (range 0.72-0.92). Meisel et al. (2020) used EDA, PPG, temperature, and ACC to achieve 51.2% sensitivity with IoC 14.1% in LOSO validation. AUC-ROC plateaus around 0.74-0.75 regardless of modality combination in forecasting applications.

### 5.4.2 Architecture Patterns

Deep learning architectures differ substantially between detection and forecasting applications. While forecasting studies concentrate on LSTM-based approaches, detection research is diverser but CNN-based approaches dominate.

Spahr et al. (2025) implemented an ensemble of 30 CNN models, achieving 96% sensitivity with 0.0054/h FPR. Elemam et al. (2025) used separate CNNs for HRV and audio processing. LSTM appears in one detection study, with Wang et al. (2025) using an LSTM with 40 hidden units processing attitude angle features. Traditional neural networks appear in earlier studies, with Fine (2025) using an ANN with 594 handcrafted features.

Three of five forecasting studies use LSTM variants. Meisel et al. (2020) implemented a minimal LSTM with 10 hidden units operating directly on raw multi-modal wrist data. Nasser et al. (2021) used a deeper 4-layer LSTM with 128 hidden units per layer. Stirling et al. (2021) combined LSTM with random forest and logistic regression in an ensemble. No forecasting study uses CNN-based approaches, possibly because forecasting requires modeling temporal dependencies rather than spatial patterns.

Handcrafted features remain competitive. Three of eight detection studies use feature-based approaches with strong performance. Fine (2025) achieved 100% sensitivity using 594 handcrafted features. This suggests that domain knowledge and human involvement

still remains valuable despite the end-to-end learning trend. Window sizes vary substantially between applications. Detection studies use shorter windows ranging from 4 seconds to 5 minutes. Forecasting studies use longer windows ranging from 30 seconds to 60 minutes.

### 5.4.3 Personalization Trade-offs

Personalization strategy represents a fundamental design choice. Global models train on population data and apply to all patients while patient-specific models train on individual patient data. Five studies use global model approaches and four studies use patient-specific approaches. Vieluf et al. (2025) used a mixed approach combining population-level harmonic patterns with individual diary data.

Global models enable on-device processing without requiring individual patient training. Spahr et al. (2025) achieved 112 ms inference time suitable for real-time processing. However, global models may not capture individual seizure patterns. Meisel et al. (2020) used LOSO validation to assess global model generalizability, achieving 62% patient success (43 of 69 patients above chance).

Patient-specific models achieve higher individual success rates but require substantial individual data for initial training. Stirling et al. (2021) achieved 100% patient success for hourly forecasting with weekly retraining. Nasser et al. (2021) achieved 83% patient success (5 of 6 patients). These rates are substantially higher than the 62% achieved by global models. However, patient-specific approaches require a mean of 14.6 months of data per patient or 60+ days per patient. This data requirement creates a cold start problem where new patients must wait weeks or months before the system becomes useful. The trade-off is clear. Patient-specific models achieve higher individual success (83-100%) compared to global models with LOSO validation (62%), but they require weeks to months of individual data collection. Global models enable immediate deployment and on-device processing but may not work well for all patients. Mixed approaches may offer a middle path, though evidence remains limited to a single study.

## 5.5 Personalization Strategies

Personalization strategy represents a fundamental design choice in wearable seizure monitoring systems. Global models train on population data and apply to all patients. Patient-specific models train on individual patient data. Mixed approaches combine population patterns with individual adaptation.

### 5.5.1 Global Models

Five studies use global model approaches. Spahr et al. (2025) trained a global model on 384 patients with tunable parameters. Fine (2025) trained on 15 patients and tested on 3 patients. Dong et al. (2022) trained on 68 patients with 10-fold cross-validation. Elemam et al. (2025) trained on 198 questionnaire responses and 30 HRV recordings. Meisel et al. (2020) trained a global model tested with LOSO validation.

Global models offer advantages for deployment. They enable on-device processing without requiring individual patient training. They can be validated with larger datasets. However, they may not capture individual seizure patterns.

Meisel et al. (2020) used LOSO validation to assess global model generalizability. The study achieved 62% patient success (43 of 69 patients above chance). This result indicates

that global models do not work equally well for all patients.

### 5.5.2 Patient-Specific Models

Four studies use patient-specific approaches. Stirling et al. (2021) trained individual models with weekly retraining for 11 patients. Nasser et al. (2021) trained separate models for each of 6 patients. Ode et al. (2023) trained anomaly detection models at the 99% confidence level for each of 66 patients. Singh Rathore et al. (2024) presented a single-patient case study.

Patient-specific models achieve higher individual success rates. Stirling et al. (2021) achieved 100% patient success for hourly forecasting. Nasser et al. (2021) achieved 83% patient success (5 of 6 patients). These rates are substantially higher than the 62% achieved by global models with LOSO validation.

Patient-specific approaches require sufficient individual data. Stirling et al. (2021) required a mean of 14.6 months of data per patient. Nasser et al. (2021) required 60+ days per patient. This data requirement creates a cold start problem where new patients must wait weeks or months before the system becomes useful.

### 5.5.3 Mixed Approaches

Vieluf et al. (2025) used a mixed approach. The study combined population-level 24-hour harmonic patterns with individual diary data. This hybrid achieved 82% sensitivity and 67% specificity. The approach may balance the advantages of global and patient-specific methods.

### 5.5.4 Subject-Split Validation

Reintjes et al. (2025) used subject-split validation. The study trained on one set of patients and tested on different patients. This approach assesses generalization to new patients. The results showed poor generalization with FPR ranging from 0.11/h to 65.62/h depending on method and configuration.

Nasser et al. (2021) used temporal split validation. Early data from each patient served as training. Later data served as testing. This approach assesses temporal stability rather than generalization to new patients.

### 5.5.5 Comparison by Strategy

Global models achieve competitive performance in detection applications. Spahr et al. (2025) achieved 96% sensitivity with 0.0054/h FPR using a global approach. This represents the best FPR among all detection studies.

Patient-specific models achieve higher individual success in forecasting applications. Stirling et al. (2021) achieved 100% patient success compared to 62% for the global model tested by LOSO in Meisel et al. (2020).

**Trade-off identified.** Patient-specific models achieve higher individual success but require weeks to months of individual data. Global models enable immediate deployment but may not work well for all patients. Mixed approaches may offer a middle path.

### 5.5.6 Computational Considerations

Global models enable on-device deployment. Spahr et al. (2025) achieved 112 ms inference time suitable for real-time processing. On-device processing enables continuous monitoring without cloud connectivity.

Patient-specific models often require cloud processing. Stirling et al. (2021) used smartphone-based processing with weekly retraining. Nasser et al. (2021) used cloud-based processing. Cloud-based approaches may experience connectivity issues and latency.

**Key limitation.** Detection studies rarely report patient-level success. Only two of eight detection studies provide individual patient results. This reporting gap prevents assessment of which patients benefit from global detection models.

## 5.6 Performance Metrics

Detection and forecasting studies use different evaluation metrics reflecting their distinct clinical objectives. Detection studies emphasize sensitivity and false positive rate (FPR). Forecasting studies report AUC-ROC, improvement over chance (IoC), and time in warning (TiW).

### 5.6.1 Detection Performance

Detection sensitivity ranges from 38.0% to 100% across studies (Table 1). Fine (2025) reported 100% sensitivity in an independent test set of 10 tonic seizures from 3 patients. Spahr et al. (2025) achieved 96% sensitivity for generalized convulsive seizures in a 384-patient multi-center study. Singh Rathore et al. (2024) reported 96.8% sensitivity in a single-patient case study. Elemam et al. (2025) achieved 95.1% sensitivity using HRV-based detection.

Reintjes et al. (2025) evaluated three anomaly detection methods (Matrix Profile, MADRID, TimeVQVAE-AD), each optimized for minimal false alarms, maximal sensitivity, or a balanced harmonic mean score. Across all methods, FAR-optimized configurations achieved false alarm rates as low as 0.06/h but missed 33%-93% of seizures, while sensitivity-optimized configurations detected 65%-98% of seizures at the cost of 4-40 false alarms per hour. This consistent pattern across methods demonstrates a fundamental sensitivity-FPR tradeoff where clinically acceptable false alarm rates below 0.1/h require sacrificing detection of most seizures, while high-sensitivity operating points generate false alarm burdens caretakers in home deployment.

Only six of nine detection studies report FPR. The ILAE Phase 3 benchmark for clinically acceptable FPR in home settings is below 0.05 to 0.1 per hour. Two studies meet this benchmark. Spahr et al. (2025) reported 0.0054/h FPR (approximately 1 false alarm per 8 days), which is approximately 10 times below the clinical benchmark. Fine (2025) reported 0.023/h FPR, approximately twice below the benchmark.

Three studies exceed the ILAE Phase 3 benchmark. Dong et al. (2022) reported 0.165/h FPR. Wang et al. (2025) reported 0.364/h FPR. Ode et al. (2023) reported 0.85/h FPR, which is 8.5 times above the benchmark. Reintjes et al. (2025) reported FPR ranging from 0.11/h to 65.62/h depending on method and configuration, which is clinically unacceptable for home use. The lowest FPR of 0.11/h was achieved by MADRID in the FAR-optimized configuration but detected only 2.44% of seizures, demonstrating the severe sensitivity-FPR tradeoff.

Three detection studies do not report FPR. Singh Rathore et al. (2024), Elemam et al. (2025), and Borujeny et al. (2013) lack standardized FPR reporting. This renders clinical assessment and comparison across studies impossible.

### 5.6.2 Forecasting Performance

Forecasting studies use different metrics than detection studies. AUC-ROC serves as the primary metric in two forecasting studies. Nasser et al. (2021) reported AUC 0.75 (SD 0.15) across 6 patients with temporal split validation. Stirling et al. (2021) reported AUC 0.74 with 100% patient success.

Sensitivity in forecasting ranges from 51.2% to 82%. Vieluf et al. (2025) reported 82% sensitivity with 67% specificity. Meisel et al. (2020) reported 51.2% sensitivity with LOSO validation.

Forecasting studies also report forecasting-specific metrics. Meisel et al. (2020) reported IoC of 14.1% and TiW of 43.7%. Nasser et al. (2021) reported TiW ranging from 0.9 to 7.2 h/day. Stirling et al. (2021) reported mean prediction time of 37 min for hourly forecasts and 3 days for daily forecasts.

Forecasting AUC consistently falls between 0.74 and 0.75 across studies. This consistency suggests a performance ceiling for non-invasive forecasting approaches using wearable biosignals.

## 5.7 Detection Versus Forecasting

Detection and forecasting serve different clinical purposes. Detection aims to identify seizures as they occur for timely intervention. Forecasting aims to predict seizures before onset to allow preventive action. These different objectives lead to distinct technical approaches and validation requirements.

### 5.7.1 Clinical Objectives

Detection systems alert caregivers when a seizure occurs. This enables intervention during or immediately after the seizure. Detection is most valuable for generalized tonic-clonic or focal to bilateral tonic-clonic seizures where caregiver response can prevent injury or complications.

Forecasting systems provide advance warning of increased seizure likelihood. This allows patients to modify activities, take rescue medication, or avoid risky situations. Forecasting could improve quality of life by reducing the uncertainty of when seizures might occur.

### 5.7.2 Algorithmic Differences

Detection and forecasting employ different algorithmic paradigms. Detection research uses diverse approaches including feature-based methods (3 studies), ensemble CNN (1 study), two-stage architectures (1 study), and anomaly detection (2 studies). Forecasting research concentrates on LSTM-based approaches (3 studies) with some use of hybrid architectures.

Detection architectures prioritize spatial pattern recognition. CNN-based approaches appear in two detection studies but no forecasting studies. Forecasting architectures prioritize temporal dependency modeling. LSTM or LSTM-ensemble approaches appear in three of five forecasting studies.

### 5.7.3 Modality Preferences

Detection and forecasting favor different biosignal modalities. Detection studies emphasize movement sensors. Six of eight detection studies use accelerometer data. Three use gyroscope data. Movement sensors capture the motor manifestations of convulsive seizures.

Forecasting studies emphasize autonomic nervous system indicators. Four of five forecasting studies use EDA. Three use ECG or HRV data. Two use temperature data. Autonomic changes may precede seizure onset by minutes to hours.

### 5.7.4 Performance Comparison

Detection and forecasting cannot be directly compared because they use different metrics. Detection sensitivity ranges from 38% to 100%. Two forecasting studies report AUC-ROC of 0.74 to 0.75, while other forecasting studies use different metrics (IoC, AUC-PR). These metrics assess different constructs.

Detection faces the FPR challenge. Only two of eight detection studies meet the clinical FPR benchmark of below 0.05 to 0.1 per hour. High FPR leads to alarm fatigue and limits real-world utility. Forecasting does not have an established clinical benchmark for acceptable performance.

### 5.7.5 Clinical Maturity

Detection shows greater clinical maturity than forecasting. Two detection studies meet Phase 3 FPR benchmarks. Commercial devices exist for detection including the Embrace watch and NightWatch armband. No forecasting study meets a clear clinical benchmark. No device has regulatory approval for seizure forecasting.

Detection benefits from clearer clinical requirements. The ILAE has defined guidelines for detection device validation including FPR targets. Forecasting lacks equivalent standardized guidelines. This gap complicates assessment of forecasting readiness.

### 5.7.6 Sample Size Disparity

Detection studies enroll larger samples than forecasting studies. Detection studies have a median sample size of 44 participants. Forecasting studies have a median sample size of 11 participants. This disparity may reflect different validation requirements.

Larger detection studies enable more robust performance estimates. Spahr et al. (2025) enrolled 384 patients across eight centers. Dong et al. (2022) enrolled 68 patients with home monitoring. The largest forecasting study, Vieluf et al. (2025), enrolled 70 patients.

### 5.7.7 Validation Setting

Detection and forecasting differ in validation settings. Six of eight detection studies were conducted in hospital or EMU settings. Only three detection studies included home validation. Three of five forecasting studies were conducted in home settings.

Home validation may be particularly important for forecasting. Forecasting systems must learn from normal daily activities to avoid false alarms from exercise, stress, or other confounding factors. Hospital settings may not capture this diversity of activities.

Nasseri et al. (2021) achieved the most rigorous validation with concurrent invasive EEG confirmation in ambulatory settings. However, only six patients were studied.

### 5.7.8 Reporting Practices

Detection and forecasting differ in metric reporting. Forecasting studies consistently report patient-level success rates. Detection studies rarely report individual patient outcomes. This reporting gap complicates assessment of which patients benefit from each approach.

Forecasting studies report forecasting-specific metrics including IoC and TiW. Detection studies emphasize sensitivity and FPR. These different metrics reflect different clinical priorities but complicate cross-domain comparison.

**Key finding.** Detection is closer to clinical translation than forecasting. Two detection approaches meet clinical FPR benchmarks. Forecasting shows consistent AUC around 0.75 but lacks clear clinical benchmarks and regulatory pathways.

### 5.7.9 Performance Ceiling

Forecasting shows a consistent performance ceiling across studies. Two forecasting studies using AUC-ROC report values of 0.74 to 0.75 (Nasseri et al. (2021), Stirling et al. (2021)). This consistency across different methods and modalities suggests a fundamental limit of non-invasive forecasting approaches. Vieluf et al. (2025) reports AUC-PR 0.84 rather than AUC-ROC, preventing direct comparison.

Detection has reached a sensitivity ceiling of 100% but struggles with FPR. The best performing detection study, Spahr et al. (2025), achieves 96% sensitivity with 0.0054/h FPR. However, most detection studies report FPR above clinical acceptability thresholds.

**Limitation identified.** Direct comparison between detection and forecasting is challenging due to different metrics, validation approaches, and clinical objectives. The field lacks standardized frameworks for assessing relative clinical value.

## 5.8 Clinical Readiness and Deployment

Clinical readiness for wearable seizure monitoring devices depends on validation phase, performance against benchmarks, and real-world deployment capability. This section assesses the translation potential of current approaches.

### 5.8.1 Validation Phase Distribution

No study in this review has achieved Phase 3 prospective validation in home settings. Phase 3 validation requires prospective testing in the intended use environment with sufficient sample size and duration.

Two detection studies approach Phase 3 requirements. Spahr et al. (2025) conducted prospective multi-center validation with 384 patients. The study met clinical FPR benchmarks but was conducted in EMU settings rather than home. Dong et al. (2022) conducted prospective home monitoring with 68 patients but reported FPR above clinical benchmarks.

Fine (2025) represents Phase 1 validation with a small independent test set. The study achieved 100% sensitivity with 0.023/h FPR but tested only 10 seizures from 3 patients. Phase 2 or 3 validation with larger samples is needed.

Forecasting studies remain in earlier validation phases. Vieluf et al. (2025) used Phase IV clinical trial data but employed retrospective analysis. Meisel et al. (2020) and Nasseri et al. (2021) represent Phase 2-equivalent validation with moderate samples.

### 5.8.2 Benchmark Achievement

The ILAE Phase 3 benchmark requires FPR below 0.05 to 0.1 per hour for home use. Only two of eight detection studies meet this benchmark.

Spahr et al. (2025) reported 0.0054/h FPR, which is approximately 10 times below the benchmark. Fine (2025) reported 0.023/h FPR, approximately twice below the benchmark. These studies demonstrate that acceptable FPR is achievable with wearable ACC-based approaches.

Six detection studies exceed the clinical FPR benchmark. Dong et al. (2022) reported 0.165/h FPR. Wang et al. (2025) reported 0.354/h FPR. Ode et al. (2023) reported 0.85/h FPR. Reintjes et al. (2025) reported FPR from 0.11/h to 65.62/h depending on method and configuration.

Forecasting lacks an established clinical benchmark. AUC around 0.74 to 0.75 appears consistently across studies, but the clinical acceptability of this performance level remains unclear. IoC of 14.1% (reported by Meisel et al. (2020)) suggests modest improvement over chance.

### 5.8.3 Commercial Device Status

Six of 13 studies use commercially available devices. The Empatica E4 research wristband appears in three studies. The Embrace watch appears in one study. The NightWatch armband appears in one study. Fitbit smartwatches appear in one study.

The Embrace watch has FDA 510(k) clearance for seizure detection. NightWatch has CE approval in Europe. However, no device has regulatory approval specifically for seizure forecasting. Forecasting devices would need to establish clinical utility and safety through dedicated trials.

Seven studies use research prototypes without clear commercialization pathways. These prototypes demonstrate technical feasibility but may require substantial engineering for commercial deployment.

### 5.8.4 Real-World Deployment

Six studies include home or ambulatory validation. Dong et al. (2022) achieved the most extensive home validation with 788 overnight recordings over three months. Stirling et al. (2021) achieved the longest monitoring duration with a mean of 14.6 months per patient. Nasser et al. (2021) achieved the most rigorous ambulatory validation with concurrent invasive EEG confirmation via the RNS system. However, only six patients were studied. Seven studies were conducted entirely in hospital or EMU settings. Hospital settings may not capture the diversity of real-world activities that cause false alarms. Exercise, cooking, household chores, and other activities present challenges not encountered in controlled hospital environments.

### 5.8.5 Deployment Architecture

Three studies report on-device or real-time deployment capability. Spahr et al. (2025) reported 112 ms inference time suitable for real-time processing. Borujeny et al. (2013) reported 316 mW power consumption for server-based processing. Dong et al. (2022) deployed on the NightWatch armband with real-time processing.

Cloud-based processing appears in three studies. Nasseri et al. (2021) and Ode et al. (2023) used cloud-based approaches. Cloud processing may introduce latency and connectivity issues. Offline processing appears in two studies without clear deployment pathways.

### 5.8.6 Barriers to Clinical Adoption

Several barriers impede clinical translation of wearable seizure monitoring. First, high FPR limits clinical utility. Six of eight detection studies report FPR above acceptable thresholds. High FPR causes alarm fatigue and may lead patients to abandon devices. Second, limited seizure type coverage restricts applicability. Most approaches detect only convulsive seizures with motor manifestations. Focal aware seizures without motor signs remain undetectable with current approaches. Non-convulsive seizures represent a substantial unmet need.

Third, small sample sizes limit generalizability. Four studies have sample sizes below 20 patients. Singh Rathore et al. (2024) is a single-patient case study. Nasseri et al. (2021) studied only 6 patients. Small samples increase uncertainty about performance estimates. Fourth, battery life and device burden limit adherence. Multi-modal systems sampling multiple signals at high frequencies consume substantial power. Long-term monitoring requires daily charging which may reduce adherence.

Fifth, data quality varies in real-world settings. Motion artifacts, skin contact problems, and environmental interference affect signal quality. Hospital validation may not capture these challenges.

### 5.8.7 Most Advanced Approaches

Spahr et al. (2025) represents the most clinically advanced detection approach. The study combines large sample size (384 patients), prospective multi-center design, excellent sensitivity (96%), and the lowest reported FPR among detection studies (0.0054/h). The algorithm operates in real-time with on-device capability. The main limitation is EMU setting rather than home validation.

Dong et al. (2022) represents the most advanced home-validated detection approach. The study achieved prospective home monitoring with the NightWatch commercial device. FPR of 0.165/h exceeds clinical benchmarks but may be acceptable for some patients.

Stirling et al. (2021) represents the most clinically advanced forecasting approach. The study achieved 100% patient success with 14.6-month home validation using consumer devices. However, forecasting lacks clear clinical benchmarks and regulatory pathways.

**Key limitation.** No study has achieved all requirements for clinical deployment. All studies have limitations in sample size, validation setting, FPR, or generalizability. Substantial additional research is needed before wearable seizure monitoring can achieve widespread clinical adoption.

## 5.9 Detection, Forecasting, and Clinical Readiness

Detection and forecasting serve different clinical purposes. Detection aims to identify seizures as they occur for timely intervention, while forecasting aims to predict seizures before onset to allow preventive action. These different objectives lead to distinct technical approaches, validation requirements, and clinical readiness levels.

### 5.9.1 Maturity Comparison

Detection shows greater clinical maturity than forecasting. Two detection studies meet Phase 3 FPR benchmarks defined by the ILAE, though both are Phase 2 studies themselves. Spahr et al. (2025) achieved 0.0054/h FPR (approximately 1 false alarm per 8 days) with 96% sensitivity in a prospective multi-center study of 384 patients. In a commentary on Larsen et al. 2024, Fine (2025) described 0.023/h FPR with 100% sensitivity on an independent test set of 3 patients with 10 tonic seizures. Commercial devices exist for detection including the Embrace watch with FDA 510(k) clearance and NightWatch with CE approval in Europe.

Forecasting shows consistent AUC around 0.74 to 0.75 across studies, suggesting a performance ceiling for non-invasive approaches. Meisel et al. (2020) reported AUC 0.74 with IoC of 14.1% over chance. Nasser et al. (2021) reported AUC 0.80 (mean) but studied only 6 patients. Vieluf et al. (2025) achieved AUC 0.75 in a larger sample of 70 patients. However, no forecasting study meets a clear clinical benchmark and no device has regulatory approval for seizure forecasting.

Detection benefits from clearer clinical requirements. The ILAE has defined guidelines for detection device validation including FPR targets below 0.05 to 0.1 per hour. Forecasting lacks equivalent standardized guidelines. This gap complicates assessment of forecasting readiness and clinical utility.

Sample size disparities further reflect the maturity difference. Detection studies have a median sample size of 66 participants compared to 40 for forecasting studies. Larger detection studies enable more robust performance estimates. The largest detection study enrolled 384 patients across eight centers, while the largest forecasting study enrolled 70 patients.

### 5.9.2 Barriers to Clinical Adoption

Several barriers impede clinical translation of wearable seizure monitoring. First, high FPR limits clinical utility. Six of eight detection studies report FPR above acceptable thresholds. Dong et al. (2022) reported 0.165/h FPR in home validation, Wang et al. (2025) reported 0.354/h, and Reintjes et al. (2025) reported FPR from 0.11/h to 65.62/h depending on method and configuration. High FPR causes alarm fatigue and may lead patients to abandon devices.

Second, limited seizure type coverage restricts applicability. Detection approaches focus on convulsive seizures with motor manifestations captured by accelerometer and gyroscope sensors. Focal aware seizures without motor signs remain undetectable with current wearable approaches. Non-convulsive seizures represent a substantial unmet need that may require invasive EEG monitoring.

Third, small sample sizes limit generalizability. Four studies have sample sizes below 20 patients. Singh Rathore et al. (2024) is a single-patient case study. Borujeny et al. (2013) studied only 3 patients. Small samples increase uncertainty about performance estimates and may not capture the full diversity of seizure presentations.

Fourth, validation setting affects real-world applicability. Seven studies were conducted entirely in hospital or EMU settings. Hospital environments may not capture the diversity of real-world activities that cause false alarms. Exercise, cooking, household chores, and other daily activities present challenges not encountered in controlled environments.

Fifth, device burden and battery life limit adherence. Multi-modal systems sampling multiple signals at high frequencies consume substantial power. Long-term monitoring

requires daily charging, which may reduce adherence over time. Stirling et al. (2021) achieved the longest monitoring duration with a mean of 14.6 months per patient, but such extended use requires careful attention to device comfort and reliability.

### 5.9.3 Most Advanced Approaches

Spahr et al. (2025) represents the most clinically advanced detection approach. The study combines large sample size, prospective multi-center design, excellent sensitivity, and the lowest reported FPR among detection studies. The ensemble 1D CNN algorithm operates in real-time with 112 ms inference time suitable for on-device processing. The main limitation is EMU setting rather than home validation.

Dong et al. (2022) represents the most advanced home-validated detection approach. The study achieved prospective home monitoring with 788 overnight recordings over three months using the NightWatch commercial device. FPR of 0.165/h exceeds clinical benchmarks but may be acceptable for some patients prioritizing sensitivity over specificity.

Stirling et al. (2021) represents the most clinically advanced forecasting approach. The study achieved 100% patient success with long home validation using consumer Fitbit devices. The LSTM+RF+LR ensemble with weekly retraining adapts to individual patient patterns. However, forecasting lacks clear clinical benchmarks and regulatory pathways.

**Key limitation.** No study has achieved all requirements for clinical deployment. All studies have limitations in sample size, validation setting, FPR, or generalizability. No study has achieved Phase 3 prospective validation in home settings with sufficient sample size and duration. Substantial additional research is needed before wearable seizure monitoring can achieve widespread clinical adoption.

## 6 Discussion

This discussion synthesizes findings from 13 studies on wearable seizure detection and forecasting. The evidence reveals different maturity levels for detection versus forecasting, distinct technical requirements, and persistent barriers to clinical translation. We organize the discussion into four parts. First, we examine clinical implications regarding device readiness and coverage gaps. Second, we analyze technical trade-offs in modality selection, architecture design, and personalization strategies. Third, we identify limitations in the evidence base that constrain confidence in the findings. Fourth, we propose research priorities to advance the field toward clinical translation.

### 6.1 Clinical Implications

From a clinical perspective, three key findings emerge regarding the scope of seizure coverage, the readiness of detection devices, and the current state of forecasting technology.

#### 6.1.1 Seizure Type Coverage Gap

Current wearable approaches address only a subset of seizure types. Accelerometer-based systems detect motor manifestations but cannot identify seizures without movement components. The ECG-based approach by Reintjes et al. (2025) examined whether autonomic signals could detect non-motor seizures. However, the observed sensitivity-FPR tradeoff (2.44% to 82.8% sensitivity at FPR ranging from 0.11/h to 65.62/h) prevents clinical

interpretation. The lowest false alarm rate detected fewer than 3% of seizures, while configurations achieving acceptable sensitivity generated false alarm burdens orders of magnitude above clinically acceptable thresholds. This fundamental tradeoff demonstrates that single-lead ECG anomaly detection cannot currently provide clinically reliable seizure detection.

Forecasting approaches may theoretically address this gap by relying on autonomic modalities rather than movement sensors. However, the consistent AUC ceiling around 0.75 suggests fundamental limits in predicting seizures without EEG biomarkers. Approximately 30% of people with epilepsy have seizures that would not be detected by current wearable approaches. This coverage gap represents a substantial unmet need that may require invasive monitoring or alternative biomarkers.

### **6.1.2 Detection Readiness**

The detection literature demonstrates that clinically acceptable FPR is achievable with accelerometer-based approaches (see Results). Two studies meet the ILAE Phase 3 benchmark using ACC data, suggesting that motor seizure detection has reached sufficient maturity for clinical translation. However, this achievement remains concentrated in controlled EMU settings rather than real-world home environments.

The EMU-to-home translation gap represents a critical barrier. The two studies achieving benchmark FPR were validated in hospital settings where patient activities are restricted and controlled. Home environments introduce motion patterns from daily living that do not occur in EMUs, including exercise, household chores, and other routine activities. The substantial FPR increase observed in home validation (0.165/h versus 0.0054/h in EMU settings) suggests that algorithm robustness to real-world activities must become a research priority.

### **6.1.3 Forecasting Status**

The forecasting literature shows a consistent performance ceiling around AUC 0.74 to 0.75 across different methods, modalities, and patient populations (see Results). This convergence suggests a fundamental limit on pre-ictal predictability using non-invasive wearable sensors. Unlike detection, where clear clinical benchmarks exist, forecasting lacks standardized validation guidelines. Without established performance targets, the clinical utility of AUC 0.75 remains uncertain.

The personalization strategy fundamentally affects forecasting utility. Global LOSO validation achieves only 43% patient success, meaning more than half of patients receive no benefit from population-based forecasting. Patient-specific approaches achieve 100% success but require weeks to months of individual data collection before the system becomes useful. This cold-start problem represents a substantial barrier to clinical adoption, as newly diagnosed patients cannot wait months for protection.

No device has regulatory approval specifically for seizure forecasting. Forecasting devices would need to establish clinical utility and safety through dedicated trials. The lack of clear benchmarks complicates pathway to regulatory approval.

## **6.2 Technical Insights and Trade-offs**

These clinical observations reflect underlying technical trade-offs in modality selection, architecture design, and personalization strategy that shape what is achievable with current

wearable approaches.

### 6.2.1 Modality Paradox

The evidence reveals a paradox regarding modality selection (see Results). Multi-modal approaches dominate the research landscape, with 69% of studies using multiple sensor types. However, the study achieving the lowest FPR uses single-modality accelerometry. This finding suggests that sensor fusion is not always necessary for reliable detection, and that biological specificity of motor manifestations may be more important than modality count.

The physiological requirements differ between detection and forecasting. Detection can succeed with movement sensors alone because ictal motor manifestations produce clear signals in ACC recordings. Forecasting requires autonomic modalities because pre-ictal changes must be detected before clinical symptoms appear. All forecasting studies incorporate EDA, ECG, or temperature data to capture these autonomic precursors. The contrast suggests that modality selection should be guided by the temporal phase of seizure activity being targeted.

### 6.2.2 Architecture Selection

Detection and forecasting applications employ different architectural patterns (see Results). Detection research shows methodological diversity including CNNs, LSTMs, ANNs, ensemble methods, and anomaly detection. This diversity may reflect the broader exploration phase of detection research. Forecasting studies concentrate heavily on LSTM variants, with three of four studies using this architecture. This convergence reflects the different temporal demands, as forecasting requires modeling longer-term dependencies than detection.

The persistence of competitive handcrafted feature approaches is noteworthy. Domain knowledge remains valuable despite the end-to-end learning trend, with engineered features achieving excellent performance in some studies. This suggests that purely data-driven approaches may not fully exploit the known physiological signatures of seizures. Ensemble methods achieve excellent results but require greater computational resources. The trade-off between ensemble complexity and deployability becomes important for commercial translation, as multi-model ensembles may be impractical for low-power wearable devices.

### 6.2.3 Personalization Dilemma

The personalization strategy presents a fundamental dilemma for clinical deployment (see Results). Global models enable immediate deployment without individual patient data, but only 43% of patients achieve performance above chance in LOSO validation. This means most patients receive no meaningful protection from population-based systems. Patient-specific approaches achieve 100% success rates but require substantial individual data collection. The cold-start problem means newly diagnosed patients must wait weeks or months before the system becomes useful. Mixed approaches combining population-level patterns with individual data may offer a middle path, but evidence remains limited to single studies.

### 6.2.4 The False Positive Challenge

The false positive challenge differs substantially between modalities (see Results). Accelerometer-based approaches can achieve clinically acceptable FPR below 0.05/h, but ECG-based approaches produce FPR an order of magnitude higher. Even the best ECG-based method exceeded the clinical benchmark by a factor of 19. This suggests that cardiac signals alone cannot reliably distinguish seizures from normal daily activities.

Multi-modal sensor fusion does not automatically solve the false positive problem. Studies combining multiple modalities still report FPR above clinical targets. The challenge may stem from the fundamental similarity between seizure-related physiological changes and those occurring during normal activities such as exercise, stress, or sleep transitions.

Forecasting faces a different performance ceiling. The consistency of AUC around 0.75 across different methods and modalities suggests fundamental limits on pre-ictal predictability using non-invasive sensors. This ceiling may reflect genuine biological variability in pre-ictal states that cannot be resolved with current wearable modalities.

## 6.3 Limitations of the Evidence Base

The strength of these conclusions is limited by evidence quality across four dimensions: sample size constraints, validation setting bias, inconsistent metrics reporting, and insufficient monitoring durations.

### 6.3.1 Sample Size Constraints

The evidence base is constrained by small sample sizes in multiple studies (see Results). Five studies have sample sizes below 20 participants, including single-patient case studies and 3-patient evaluations. Small samples limit the reliability of performance estimates and reduce confidence that results will generalize to broader epilepsy populations. The forecasting literature is particularly affected, with studies limited to 6-11 patients due to invasive EEG confirmation requirements.

### 6.3.2 Validation Setting Bias

Eight of 13 studies were conducted in hospital or EMU settings, creating a systematic bias that may underestimate real-world FPR. Hospital environments lack the diversity of daily activities that generate false alarms in home settings. The substantial FPR difference between EMU validation (0.0054/h) and home validation (0.165/h) demonstrates that hospital-based results may not translate to real-world use.

### 6.3.3 Reporting Inconsistency

Inconsistent metrics reporting hinders cross-study comparison (see Results). Forecasting studies routinely report patient-level success rates, enabling assessment of which patients benefit from the technology. Detection studies rarely provide this granularity, with only two of nine studies reporting individual patient outcomes. This reporting gap prevents assessment of whether detection algorithms work reliably for all patients or only a subset. High aggregate sensitivity may mask substantial variability across patients that remains invisible without patient-level reporting.

### 6.3.4 Monitoring Duration and Cold-Start Problems

Patient-specific approaches require weeks to months of individual data before becoming useful, creating a cold-start problem for newly diagnosed patients. This requirement conflicts with the clinical need for immediate protection after diagnosis. While global models enable immediate deployment, they achieve substantially lower patient success rates. The trade-off between immediate deployment and individualized performance remains unresolved without evidence from hybrid approaches.

## 6.4 Future Directions

Building on these limitations, nine research priorities emerge to advance wearable seizure monitoring toward clinical utility.

First, forecasting needs standardized clinical benchmarks. Detection has the ILAE Phase 3 benchmark of FPR below 0.05 to 0.1 per hour. Forecasting lacks an equivalent standard. AUC values around 0.74 to 0.75 appear consistently across forecasting studies, but the clinical meaning of this performance level remains unclear. Metrics such as IoC and TiW need standardization. The field should establish minimum acceptable performance levels for forecasting devices to enable clinical translation.

Second, large-scale home validation is needed. Only two detection studies include prospective home validation. Hospital and EMU settings may not capture the diversity of real-world activities that cause false alarms. Exercise, cooking, household chores, and other daily activities present challenges not encountered in controlled environments. The 0.165/h FPR reported by Dong et al. (2022) in home settings exceeds the clinical benchmark. Future studies should prioritize prospective validation in intended use environments with sufficient sample sizes and duration.

Third, novel modalities are needed for non-convulsive seizures. Current approaches detect only motor seizures with convulsive manifestations. Focal aware seizures without motor signs remain undetectable with current wearable approaches. Spahr et al. (2025) achieved excellent performance for convulsive seizures using accelerometer data alone, but this approach cannot detect seizures without movement. Multi-modal autonomic approaches show promise but have not addressed non-convulsive seizure types. Novel biosignals or advanced feature engineering may be required to detect the subtle physiological changes associated with focal seizures.

Fourth, adaptive personalization strategies could address the trade-off between deployability and individual performance. Global models enable immediate deployment but may not work equally well for all patients. Meisel et al. (2020) achieved 43% patient success with LOSO validation, indicating substantial inter-patient variability. Patient-specific models achieve higher individual success but require weeks to months of individual data. Stirling et al. (2021) required a mean of 14.6 months per patient. Nasseri et al. (2021) required 6+ months per patient (median 220 days). Mixed approaches that combine population-level patterns with individual adaptation, such as the method used by Vieluf et al. (2025), may offer a balance. Research should address the cold-start problem by developing models that can initialize from population data and adapt efficiently to individual patients.

Fifth, real-time processing and edge deployment require more attention. Most studies in this review ran algorithms offline on PCs or servers rather than on wearable devices. Clinical utility demands real-time detection and forecasting on-device to enable timely interventions. This requires optimization for low power consumption and limited memory. Spahr et al. (2025) is an exception, achieving 112 ms inference time suitable for smartwatch

integration. Future work should report computational requirements and demonstrate feasible deployment on target hardware.

Sixth, model explainability and interpretability are needed for clinical adoption. Deep learning architectures such as CNNs, LSTMs, and autoencoders operate as black boxes with limited transparency. Clinicians and patients need to understand what features trigger detections or forecasts to trust these systems. Nasser et al. (2021) performed ablation studies to examine feature contributions, and Reintjes et al. (2025) analyzed feature importance. However, systematic interpretability frameworks are missing. Future research should integrate explainable AI techniques and provide interpretable outputs for clinical decision-making.

Seventh, regulatory pathways and clinical integration need consideration. None of the reviewed studies address FDA or CE marking requirements for medical devices. Integration with clinical workflows, liability considerations for failed detections or forecasts, and ethical implications remain unexplored. Device approval pathways require rigorous validation data, standardized manufacturing processes, and post-market surveillance plans. Research should engage with regulatory frameworks early in development to facilitate clinical translation.

Eighth, algorithm robustness and generalizability require multi-site validation. Most studies used single-site or single-country data, limiting generalizability across populations and healthcare systems. Multi-continental validation with diverse demographic, geographic, and clinical characteristics is needed. Cross-device validation is also lacking. Studies typically evaluate algorithms on specific hardware such as Empatica E4, but performance may differ on consumer devices like Apple Watch or Fitbit. Standardized datasets and benchmarks would enable fair comparison across approaches.

Ninth, long-term algorithm stability remains unaddressed. Seizure patterns evolve due to medication changes, disease progression, and physiological changes over time. Stirling et al. (2021) retrained models weekly, acknowledging performance drift. However, no study reports systematic evaluation of model validity over periods longer than several months. Longitudinal studies assessing how detection and forecasting performance decays over time are needed. Adaptive algorithms that can detect and compensate for performance degradation would improve clinical utility.

## 7 Conclusion

This systematic review analyzed 13 studies on wearable seizure detection and forecasting using deep learning approaches. The evidence base comprises 912 participants across nine detection studies and four forecasting studies published between 2013 and 2026.

The predominance of accelerometer-based detection reveals a fundamental constraint in wearable seizure monitoring. Six of nine detection studies rely on accelerometry, and the two studies achieving ILAE Phase 3 benchmarks for false positive rates both use ACC. Spahr et al. (2025) achieved 0.0054/h FPR with single-modality ACC. Fine (2025) achieved 0.023/h FPR with a 6-axis ACC and gyroscope band. This convergence on a single modality reflects the biological reality that convulsive seizures produce distinct motor signatures that deep learning can extract reliably. However, it also means current wearable approaches can serve only the subset of patients with motor manifestations. Approximately 30% of people with epilepsy have non-convulsive seizures that remain undetectable with movement-based approaches.

Forecasting studies face a different constraint. The consistent AUC plateau around 0.74 to

0.75 across studies using different methods, modalities, and patient populations suggests a fundamental limit on pre-ictal predictability using non-invasive sensors. Unlike detection, where the ictal motor event produces a clear signal, forecasting must rely on autonomic precursors that may be inherently less specific. The autonomic nervous system responds to stress, exercise, sleep transitions, and other daily events. These confounds may establish a performance ceiling for wearable forecasting that cannot be breached without EEG biomarkers or novel sensing modalities.

No study in this review has achieved all requirements for widespread clinical adoption. Commercial detection devices exist with FDA or CE approval, but the deep learning approaches reviewed here lack large-scale prospective validation. Forecasting faces an additional barrier. No device has regulatory approval for forecasting, and the field lacks established clinical benchmarks equivalent to the ILAE Phase 3 FPR standard. The performance achieved to date, while scientifically promising, may not be sufficient for regulatory approval or clinical utility.

The clinical pathway forward requires addressing three gaps. First, home validation at scale is needed to bridge the EMU-to-home translation gap. The two studies meeting FPR benchmarks were conducted in hospital settings. Real-world deployment introduces activities and confounds not captured in EMU environments. Second, forecasting needs standardized clinical benchmarks to define what level of AUC, IoC, or TiW constitutes clinically meaningful performance. Third, novel modalities or approaches are needed for non-convulsive seizures. ECG-based detection shows some promise for identifying autonomic changes during non-motor seizures, but current performance remains insufficient. Mixed personalization strategies that combine population patterns with individual adaptation may address the deployability versus performance trade-off without requiring months of individual data collection.

## References

- AbuAlrob, M. A., Itbaisha, A., & Mesraoua, B. (2025). Unlocking new frontiers in epilepsy through AI: From seizure prediction to personalized medicine. *Epilepsy & Behavior*, 166, 110327. Retrieved November 21, 2025, from URL: <https://www.sciencedirect.com/science/article/pii/S1525505025000666> (cit. on p. 4).
- Beniczky, S., & Ryvlin, P. (2018). Standards for testing and clinical validation of seizure detection devices. *Epilepsia*, 59 Suppl 1, 9–13. DOI: [10.1111/epi.14049](https://doi.org/10.1111/epi.14049) (cit. on p. 4).
- Beniczky, S., Wiebe, S., Jeppesen, J., Tatum, W. O., Brazdil, M., Wang, Y., Herman, S. T., & Ryvlin, P. (2021). Automated seizure detection using wearable devices: A clinical practice guideline of the International League Against Epilepsy and the International Federation of Clinical Neurophysiology. *Clinical Neurophysiology*, 132(5), 1173–1184. DOI: [10.1016/j.clinph.2020.12.009](https://doi.org/10.1016/j.clinph.2020.12.009) (cit. on pp. 3–5).
- Borujeny, G. T., Yazdi, M., Keshavarz-Haddad, A., & Borujeny, A. R. (2013). Detection of Epileptic Seizure Using Wireless Sensor Networks. *Journal of Medical Signals and Sensors*, 3(2), 63–68. Retrieved January 5, 2026, from URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC3788195/> (cit. on pp. 14, 15, 17, 18, 22, 25, 27).
- Chen, F., Chen, I., Zafar, M., Sinha, S. R., & Hu, X. (2022). Seizures detection using multimodal signals: A scoping review. *Physiological Measurement*, 43(7), 07TR01. DOI: [10.1088/1361-6579/ac7a8d](https://doi.org/10.1088/1361-6579/ac7a8d) (cit. on pp. 3, 4).

- Dong, C., Ye, T., Long, X., Aarts, R. M., van Dijk, J. P., Shang, C., Liao, X., Chen, W., Lai, W., Chen, L., & Wang, Y. (2022). A Two-Layer Ensemble Method for Detecting Epileptic Seizures Using a Self-Annotation Bracelet With Motor Sensors. *IEEE Transactions on Instrumentation and Measurement*, 71, 1–13. DOI: [10.1109/TIM.2022.3173270](https://doi.org/10.1109/TIM.2022.3173270) (cit. on pp. 14–16, 19, 21, 23–28, 32).
- Elemam, M. E., Elgamal, A., Elmenshawi, I., & Abdelkader, H. E. (2025). Automated validated tool for epileptic seizure detection using deep learning. *Journal of Intelligent Systems and Internet of Things*, 15(1), 74–90. DOI: [10.54216/JISIoT.150107](https://doi.org/10.54216/JISIoT.150107) (cit. on pp. 16–19, 21, 22).
- Fine, A. L. (2025). Detection is Key: Automated Tonic Seizure Detection With a Wearable Device. *Epilepsy Currents*, 25(1), 36–38. DOI: [10.1177/15357597241293298](https://doi.org/10.1177/15357597241293298) (cit. on pp. 14–19, 21, 24, 25, 27, 33).
- Ghaderi, F. (2025, April 11). *Advances in Machine Learning for Epileptic Seizure Prediction: A Review of ECG-Based Approaches*. 2025040942. DOI: [10.20944/preprints202504.0942.v1](https://doi.org/10.20944/preprints202504.0942.v1). (Cit. on p. 4).
- Hussein, A. F., Arunkumar, N., Gomes, C., Alzubaidi, A. K., Habash, Q. A., Santamaria-Granados, L., Mendoza-Moreno, J. F., & Ramirez-Gonzalez, G. (2018). Focal and Non-Focal Epilepsy Localization: A Review. *IEEE Access*, 6, 49306–49324. DOI: [10.1109/ACCESS.2018.2867078](https://doi.org/10.1109/ACCESS.2018.2867078) (cit. on p. 3).
- Jain, S., Oswal, U., Xu, K. S., Eriksson, B., & Haupt, J. (2017). A Compressed Sensing Based Decomposition of Electrodermal Activity Signals. *IEEE Transactions on Biomedical Engineering*, 64(9), 2142–2151. DOI: [10.1109/TBME.2016.2632523](https://doi.org/10.1109/TBME.2016.2632523) (cit. on p. 4).
- Kalousios, S., Müller, J., Yang, H., Eberlein, M., Uckermann, O., Schackert, G., Polanski, W. H., & Leonhardt, G. (2024). ECG-based epileptic seizure prediction: Challenges of current data-driven models. *Epilepsia Open*, 10(1), 143–154. DOI: [10.1002/epi4.13073](https://doi.org/10.1002/epi4.13073) (cit. on p. 5).
- Leal, A., da Graça Ruano, M., Henriques, J., de Carvalho, P., & Teixeira, C. (2017). On the viability of ECG features for seizure anticipation on long-term data. *2017 IEEE 3rd International Forum on Research and Technologies for Society and Industry (RTSI)*, 1–5. DOI: [10.1109/RTSI.2017.8065951](https://doi.org/10.1109/RTSI.2017.8065951) (cit. on p. 4).
- Lu, S., Liu, L., Li, J., Chambers, J., Cook, M. J., & Grayden, D. B. (2025). Leveraging Channel Coherence in Long-Term iEEG Data for Seizure Prediction. *IEEE Journal of Biomedical and Health Informatics*, 29(8), 5541–5548. DOI: [10.1109/JBHI.2025.3556775](https://doi.org/10.1109/JBHI.2025.3556775) (cit. on p. 5).
- Mason, F., Scarabello, A., Taruffi, L., Pasini, E., Calandra-Buonaura, G., Vignatelli, L., & Bisulli, F. (2024). Heart Rate Variability as a Tool for Seizure Prediction: A Scoping Review. *Journal of Clinical Medicine*, 13(3), 747. DOI: [10.3390/jcm13030747](https://doi.org/10.3390/jcm13030747) (cit. on p. 4).
- Meisel, C., El Atrache, R., Jackson, M., Schubach, S., Ufongene, C., & Loddenkemper, T. (2020). Machine learning from wristband sensor data for wearable, noninvasive seizure forecasting. *Epilepsia*, 61(12), 2653–2666. DOI: [10.1111/epi.16719](https://doi.org/10.1111/epi.16719) (cit. on pp. 16–20, 22, 24, 25, 27, 32).
- Miron, G., Halimeh, M., Jeppesen, J., Loddenkemper, T., & Meisel, C. (2025). Autonomic biosignals, seizure detection, and forecasting. *Epilepsia*, 66 Suppl 3, 25–38. DOI: [10.1111/epi.18034](https://doi.org/10.1111/epi.18034) (cit. on pp. 4, 5).
- Nasseri, M., Pal Attia, T., Joseph, B., Gregg, N. M., Nurse, E. S., Viana, P. F., Worrell, G., Dümpelmann, M., Richardson, M. P., Freestone, D. R., & Brinkmann, B. H.

- (2021). Ambulatory seizure forecasting with a wrist-worn device using long-short term memory deep learning. *Scientific Reports*, 11(1), 21935. DOI: [10.1038/s41598-021-01449-2](https://doi.org/10.1038/s41598-021-01449-2) (cit. on pp. 16–27, 32, 33).
- Ode, R., Fujiwara, K., Miyajima, M., Yamakawa, T., Kano, M., Jin, K., Nakasato, N., Sawai, Y., Hoshida, T., Iwasaki, M., Murata, Y., Watanabe, S., Watanabe, Y., Suzuki, Y., Inaji, M., Kunii, N., Oshino, S., Khoo, H. M., Kishima, H., & Maehara, T. (2023). Development of an epileptic seizure prediction algorithm using R–R intervals with self-attentive autoencoder. *Artificial Life and Robotics*, 28(2), 403–409. DOI: [10.1007/s10015-022-00832-0](https://doi.org/10.1007/s10015-022-00832-0) (cit. on pp. 15–18, 20, 21, 25, 26).
- Pavei, J., Heinzen, R. G., Novakova, B., Walz, R., Serra, A. J., Reuber, M., Ponnusamy, A., & Marques, J. L. B. (2017). Early Seizure Detection Based on Cardiac Autonomic Regulation Dynamics. *Frontiers in Physiology*, 8. DOI: [10.3389/fphys.2017.00765](https://doi.org/10.3389/fphys.2017.00765) (cit. on p. 4).
- Reintjes, C., Hagenbeck, J. F., Ballo, M., Rahlmeier, T., Wolf, S. M., & Schoder, D. (2025). ECG-Based Detection of Epileptic Seizures in Real-World Wearable Settings: Insights from the SeizeIT2 Dataset. *Sensors*, 25(24), 7687. DOI: [10.3390/s25247687](https://doi.org/10.3390/s25247687) (cit. on pp. 15, 16, 18, 20, 21, 25, 27, 28, 33).
- Seth, E. A., Watterson, J., Xie, J., Arulsamy, A., Md Yusof, H. H., Ngadimon, I. W., Khoo, C. S., Kadirvelu, A., & Shaikh, M. F. (2023). Feasibility of cardiac-based seizure detection and prediction: A systematic review of non-invasive wearable sensor-based studies. *Epilepsia Open*, 9(1), 41–59. DOI: [10.1002/epi4.12854](https://doi.org/10.1002/epi4.12854) (cit. on p. 4).
- Shum, J., & Friedman, D. (2021). Commercially available seizure detection devices: A systematic review. *Journal of the Neurological Sciences*, 428, 117611. DOI: [10.1016/j.jns.2021.117611](https://doi.org/10.1016/j.jns.2021.117611) (cit. on p. 5).
- Singh Rathore, S. P., Magotra, S., K V, S., Sonu Kumar, S., Kaur, G., & Singh, S. (2024). Development of a Multimodal Machine Learning Model for Seizure Detection Using Wearable Devices. *2024 1st International Conference on Advances in Computing, Communication and Networking (ICAC2N)*, 1422–1426. DOI: [10.1109/ICAC2N63387.2024.10895217](https://doi.org/10.1109/ICAC2N63387.2024.10895217) (cit. on pp. 14, 16, 17, 20–22, 26, 27).
- Sivathamboo, S., Nhu, D., Piccenna, L., Yang, A., Antonic-Baker, A., Vishwanath, S., Todaro, M., Yap, L. W., Kuhlmann, L., Cheng, W., O’Brien, T. J., Lannin, N. A., & Kwan, P. (2022). Preferences and User Experiences of Wearable Devices in Epilepsy: A Systematic Review and Mixed-Methods Synthesis. *Neurology*, 99(13), e1380–e1392. DOI: [10.1212/WNL.0000000000200794](https://doi.org/10.1212/WNL.0000000000200794) (cit. on p. 4).
- Spahr, A., Bernini, A., Ducouret, P., Baumgartner, C., Koren, J. P., Imbach, L., Beniczky, S., Larsen, S. A., Rheims, S., Fabricius, M., Seeck, M., Steinhoff, B. J., Beuchat, I., Dan, J., Atienza, D. A., Bardyn, C.-E., & Ryvlin, P. (2025). Deep learning-based detection of generalized convulsive seizures using a wrist-worn accelerometer. *Epilepsia*, 66 Suppl 3, 53–63. DOI: [10.1111/epi.18406](https://doi.org/10.1111/epi.18406) (cit. on pp. 14–21, 23–28, 32, 33).
- Stirling, R. E., Grayden, D. B., D’Souza, W., Cook, M. J., Nurse, E., Freestone, D. R., Payne, D. E., Brinkmann, B. H., Pal Attia, T., Viana, P. F., Richardson, M. P., & Karoly, P. J. (2021). Forecasting Seizure Likelihood With Wearable Technology. *Frontiers in Neurology*, 12, 704060. DOI: [10.3389/fneur.2021.704060](https://doi.org/10.3389/fneur.2021.704060) (cit. on pp. 14, 16–22, 24–26, 28, 32, 33).

- Sveinsson, O., Andersson, T., Mattsson, P., Carlsson, S., & Tomson, T. (2020). Clinical risk factors in SUDEP. *Neurology*, *94*(4), e419–e429. DOI: [10.1212/WNL.00000000000008741](https://doi.org/10.1212/WNL.00000000000008741) (cit. on p. 3).
- Tugwell, P., & Tovey, D. (2021). PRISMA 2020. *Journal of Clinical Epidemiology*, *134*, A5–A6. DOI: [10.1016/j.jclinepi.2021.04.008](https://doi.org/10.1016/j.jclinepi.2021.04.008) (cit. on p. 5).
- Vieluf, S., Tomioka, S., Zhang, B., Krishnan, V., Bosl, W. J., Grinnell, T., & Loddenkemper, T. (2025). Seizure monitoring by combined diary and wearable data: A multicenter, longitudinal, observational study. *Epilepsia*, *66*(11), 4259–4271. DOI: [10.1111/epi.18550](https://doi.org/10.1111/epi.18550) (cit. on pp. 16, 17, 19, 20, 22–24, 27, 32).
- Wang, J., Hu, D., Lai, X., Jiang, T., Jiang, T., Gao, F., Vidal, P.-P., & Cao, J. (2025). Epileptic Seizure Detection Based on Attitude Angle Signal of Wearable Device. *IEEE Transactions on Instrumentation and Measurement*, *74*. DOI: [10.1109/TIM.2025.3529058](https://doi.org/10.1109/TIM.2025.3529058) (cit. on pp. 4, 15–18, 21, 25, 27).
- Webster, J., & Watson, R. T. (2002). Analyzing the past to prepare for the future: Writing a literature review. *MIS quarterly*, xiii–xxiii. Retrieved December 3, 2025, from URL: <https://www.jstor.org/stable/4132319> (cit. on pp. 2, 5, 7).
- Wong, S., Simmons, A., Rivera-Villicana, J., Barnett, S., Sivathamboo, S., Perucca, P., Ge, Z., Kwan, P., Kuhlmann, L., Vasa, R., Mouzakis, K., & O’Brien, T. J. (2023). EEG datasets for seizure detection and prediction— A review. *Epilepsia Open*, *8*(2), 252–267. DOI: [10.1002/epi4.12704](https://doi.org/10.1002/epi4.12704) (cit. on p. 5).
- Wu, D., Wei, J., Vidal, P.-P., Wang, D., Yuan, Y., Cao, J., & Jiang, T. (2024). A Novel Seizure Detection Method Based on the Feature Fusion of Multimodal Physiological Signals. *IEEE Internet of Things Journal*, *11*(16), 27545–27556. DOI: [10.1109/JIOT.2024.3398418](https://doi.org/10.1109/JIOT.2024.3398418) (cit. on p. 4).